1   Eye-tracking on the web: lessons learned from replicating 6 experiments

2   First Author[1] & Ernst-August Doelle[1,2]

3   [1] Wilhelm-Wundt-University

4   [2] Konstanz Business School

5   Author Note

6   Add complete departmental affiliations for each author here. Each new line herein

7   must be indented, like this line.

8   Enter author note here.

9   The authors made the following contributions. First Author: Conceptualization,

10   Writing - Original Draft Preparation, Writing - Review & Editing; Ernst-August Doelle:

11   Writing - Review & Editing.

12   Correspondence concerning this article should be addressed to First Author, Postal

13   address. E-mail: my@email.com

Abstract

ADD LATER

Keywords: keywords

Word count: X

18   Eye-tracking on the web: lessons learned from replicating 6 experiments

19   Intro stuff:

20   • Eye-tracking as a key method in cognitive science research

21   • Online data collection is more and more popular & let's us ask new questions

22   • But, concerns over quality + little known about eye-tracking online

23   *Present work*

24   In the present work, we attempted to replicate six eye-tracking studies from the

25 cognitive science literature using the eye-tracking plug-in from `jsPsych`, a Javascript

26 library for running behavioral experiments in a web browser (de Leeuw, 2015). The plug-in

27 relies on the `webgazer.js` library for webcam eye-tracking (Papoutsaki et al., 2016). The

28 goal was to examine the strengths and weaknesses of webcam eye-tracking for common

29 paradigms in cognitive science. The studies were chosen to cover a variety of topic areas

30 (e.g., memory, decision-making, psycholinguistics) and paradigms (two halves of the screen,

31 visual world paradigm with four quadrants, visual world paradigm with "natural" scenes).

32 . . .

33   **General Methods**

34 **Participants**

35   Participants completed the experiment remotely and were recruited through the

36 Prolific platform. In order to have access to the experiment, participants had to meet the

37 following criteria: 18 years of age or older, fluency in English, and access to a webcam. All

38 participants provided informed consent. The studies were approved by the Vassar College

39 Institutional Review Board.

Eye-tracking Calibration and Validation

When participants began the experiment, they were notified the webcam would be used for eye tracking but no video would be saved. They were asked to remove glasses if possible, close any other tabs or apps, turn off notifications, and make sure their face was lit from the front. The webcam's view of the participant popped up on the screen, and participants were asked to center their face in the box and keep their head still. The experiment window then expanded to full screen, and participants began the eye-tracking calibration.

During the calibration, dots appeared on the screen one at a time in different locations, and the participants had to fixate them and click on each one. Once they clicked on a dot, it would disappear and a new one would appear in a different location on the screen. The locations of calibration dots were specific to each experiment (details below) and appeared in the areas of the screen where the visual stimuli would appear during the main task in order to ensure that eye movements were accurately recorded in the relevant regions of interest. After the calibration was completed, the validation began. Participants were asked to go through the same steps as the calibration, except that they only fixated the dots as they appeared in different locations on the screen. If accuracy on the validation was too low (fewer than 50% of looks landed within a 200 px radius of the validation points), participants were given an opportunity to re-start the calibration and validation steps. If the second attempt also lead to low validation accuracy, participants were informed that they could not participate in the study.

Data pre-processing

We used R [Version 4.1.2; R Core Team (2021)] and the R-packages *afex* [Version 1.0.1; Singmann, Bolker, Westfall, Aust, and Ben-Shachar (2021)], *broom.mixed* [Version 0.2.7; Bolker and Robinson (2020)], *dplyr* [Version 1.0.7; Wickham, François, Henry, and

Müller (2021)], *forcats* [Version 0.5.1; Wickham (2021a)], *ggplot2* [Version 3.3.5; Wickham (2016)], *jsonlite* [Version 1.7.2; Ooms (2014)], *lme4* [Version 1.1.27.1; Bates, Mächler, Bolker, and Walker (2015)], *lmerTest* [Version 3.1.3; Kuznetsova, Brockhoff, and Christensen (2017)], *Matrix* [Version 1.3.4; Bates and Maechler (2021)], *papaja* [Version 0.1.0.9997; Aust and Barth (2020)], *readr* [Version 2.0.2; Wickham and Hester (2020)], *shiny* [Version 1.7.1; Chang et al. (2021)], *stringr* [Version 1.4.0; Wickham (2019)], and *tidyr* [Version 1.1.4; Wickham (2021b)] for all our analyses.

Experiment 1

The first study was a replication attempt of Altmann and Kamide (1999). Altmann and Kamide used the visual world eye-tracking paradigm (Tanenhaus, Spivey-Knowlton, Eberhard, & Sedivy, 1995) to show that meanings of verbs rapidly constrain the set of potential subsequent referents in sentence processing. For example, when looking at the display in Figure XX and listening to a sentence like "The boy will eat the...," participants are more likely to look at the cake than when they hear "The boy will move the...," in which case they tend to look at the train, presumably because cakes are edible and trains are not. Semantic information available at the verb is used to anticipate upcoming linguistic input.

Methods

All stimuli, experiment scrips, data, analysis scripts, and pre-registration are available on the Open Science Framework at https://osf.io/s82kz.

Participants. 60 participants were paid $XX for their participation. Our sample size of participants was determined by the total run time of our experiment, ~10 minutes, and the allotted funding from the Vassar College Cognitive Science Department. From this information, we calculated a reasonable number of participants we could afford to

compensate on Prolific. Note that the sample size of the original study was 24. For unknown reasons, 2 of the subjects' results were not recorded, so in the analysis, we worked with data collected from 58 participants.

Procedure. The task began with an -point eye-tracker calibration and validation. During the experiment, the participants were simultaneously presented with a visual image and a corresponding audio recording of a spoken sentence. Participants had to input a keyboard response indicating "yes" or "no" as to whether the sentence they heard was feasible given the visual image. There were two practice trials to ensure that participants had a sound understanding of the instructions before they undertook the main portion of the experiment. Participants' reaction times, keyboard responses, and looks to objects in the scene were recorded for each trial.

Materials & Design. The visual stimuli were created through Canva and depicted an agent accompanied by four to five objects in the scene (see Figure XX). On critical trials, participants heard one of two sentences associated with the scene. In the restrictive condition, the sentence (e.g., "The boy will eat the cake") contained a verb (e.g., "eat") which restricts the set of possible subsequent referents (e.g., to edible things). Only the target object (e.g., the cake) was semantically consistent with the verb's meaning. In the non-restrictive condition, the sentence (e.g., "The boy will move the cake") contained a verb (e.g., "move") which does not restrict the set of possible subsequent referents. The target object (e.g., the cake) as well as the distractor objects (e.g., the train, the ball, etc.) were semantically consistent with the verb's meaning. Both sentences were compatible with the scene, such that the correct keyboard response for the critical trials was "yes." Filler trials consisted of scenes that looked similar to critical scenes but were paired with inappropriate sentences. The correct keyboard response for the filler trials was "no."

Each participant was presented with sixteen critical trials (eight in the restrictive condition, eight in the non-restrictive condition) and sixteen fillers for a total of 32 trials. The order of trials and the assignment of critical scene to condition was random on a

subject-by-subject basis.

TO DO: add figure

Data pre-processing and analysis. Looks to the objects in the scene were time-locked to the onset of the verb, the offset of the verb, onset of the post-verbal determiner, and onset of the target noun.

Results

Replication.

- here we will describe the analyses that are as close as possible to the original paper with a minimal validation cutoff

- same analysis but with stricter validation cutoff

Comparison to in-lab data.

- here we will describe a direct comparison to data collected in the lab

Calibration.

- here we will describe the analyses that correlate calibration quality with effect size at the individual level

Discussion

<div align="center">

Experiment 2

</div>

The second study was a replication attempt of Johansson and Johansson (2014), which examined how visuospatial information is integrated into memory for objects. They found that, during memory retrieval, learners spontaneously look to blank screen locations where pictures were located during encoding (see Spivey & Geng, 2001) and that this spatial reinstatement facilitates retrieval of the picture.

Methods

Participants.   60 participants were paid $XX for their participation . The sample size was motivated in part by budget constraints, but was nonetheless 2.5x larger than the original sample size of 24, as suggested by Simonsohn (2015). Data from 1 participant were not properly recorded due to unknown technical issues, so data from 59 participants were included in all analyses to follow.

Procedure.   The experiment consisted of two blocks each composed of an encoding phase and a recall phase. During the encoding phase, participants saw a grid indicating the four quadrants of the screen. Each quadrant contained six images of items belonging to the same category (see Figure XX). The four categories were humanoids, household objects, animals, and methods of transportation. Participants were asked to remember the contents of the four quadrants. Different images were used in each block.

Each of the four quadrants was presented one at a time. First, a list of the items in the quadrant were shown, then the items in the actual quadrant were shown (??). For each item, an audio file would play ("???") asking the participant to use their arrow keys to identify which direction each item was facing (every item was facing either left or right (right??)). After the participant identified the direction of each item, they would have an additional 30 seconds to encode the name and orientation of each item in the quadrant. Then, after all four quadrants were presented in this way, the participant was shown the full grid of 24 items and had 60 seconds to further encode the name and orientation of each item.

During the recall phase, participants listened to statements and responded by pressing the 'F' key for false statements and 'T' for true ones. Each statement fell into either an interobject or intraobject condition. Interobject statements were those that compared two different items in the grid (e.g. "The skeleton is to the left of the robot"), while intraobject statements were those that asked about the orientation of a single item

164 (e.g. "The bus is facing right"). There were 48 total statements, with 24 interobject and 24

165 intraobject statements split evenly among the four quadrants. While listening to these

166 statements, in the free-viewing block, participants saw a blank screen and were allowed to

167 freely gaze around the screen. During the fixed-viewing block, participants were asked to

168 fixate a small cross in the center of the screen throughout the recall phase. In both cases,

169 the mouse was obscured from the screen. Participants were randomly assigned to see the

170 fixed-viewing or free-viewing block first.

171 After completing both encoding-recall blocks, participants were asked to answer a few

172 survey questions (such as whether they wore glasses or encountered any distractions).

173 The primary methodological difference between this replication and Johansson and

174 Johansson's study was that the original study included two additional viewing conditions

175 that were omitted from this replication due to time constraints. In those two conditions,

176 participant were prompted to look to a specific quadrant (rather than free viewing or

177 central fixation) which either matched or mismatched the original location of the

178 to-be-remembered item.

179 **Data analysis.**

180 **Results**

181 **Replication. Eye-gaze**. Looks during the retrieval period were categorized as

182 belonging to one of four quadrants based on the x,y coordinates. The critical quadrant was

183 the one in which the to-be-retrieved object had been previously located during encoding.

184 The other three quadrants were semi-randomly labeled "first," "second," third" (e.g., when

185 the critical quadrant was in the top left, the "first" quadrant was the top right quadrant,

186 but when the critical quadrant was in the top right, "first" corresponded to bottom right,

187 etc.). In the free-viewing condition, participants directed a larger proportion of looks to the

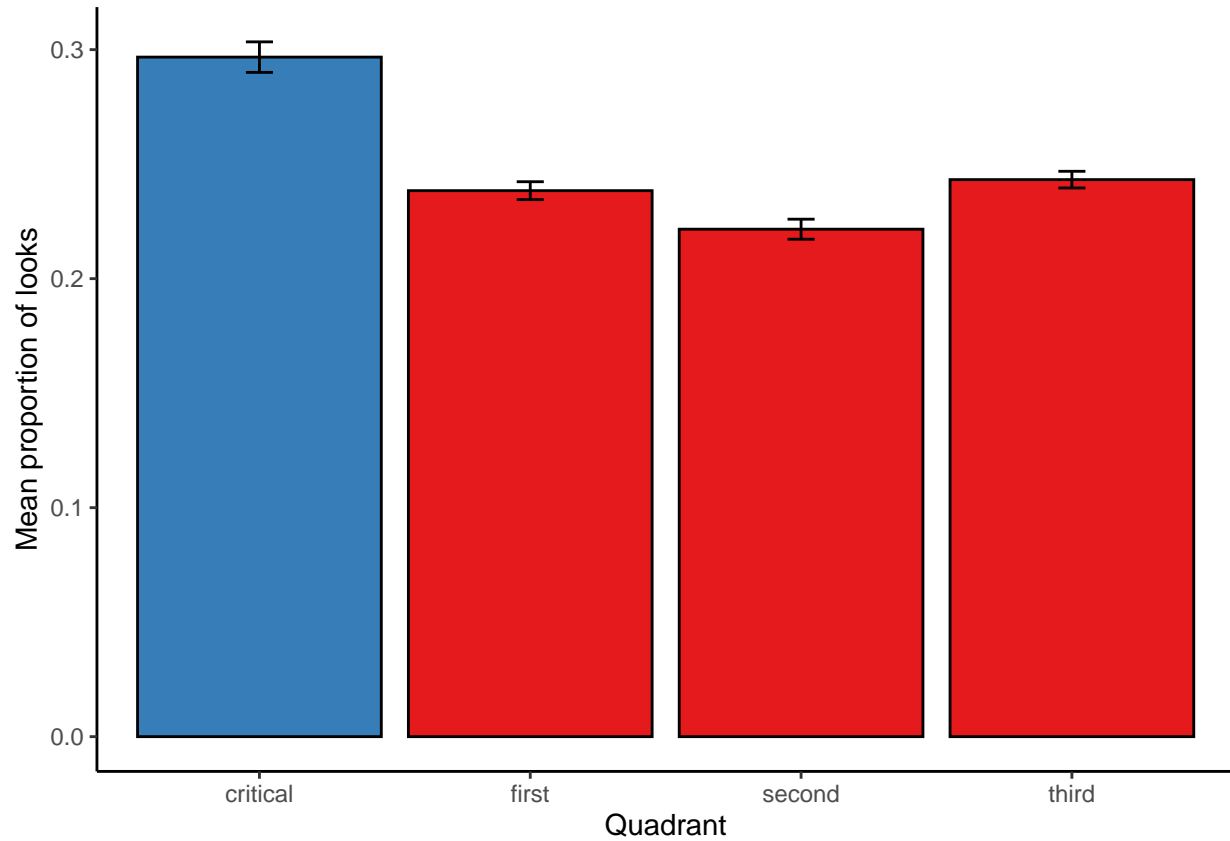188 critical quadrant (see Figure 1).

Figure 1. Proportion of eye-gaze to critical quadrant and other three quadrants durign memory retrieval.

189 The proportion of looks across quadrant was analyzed in linear mixed-effects model

190 with quadrant as the predictor (critical as the reference level). The model included random

191 intercepts and slopes for participants[1] Proportions of looks were significantly higher for the

192 critical quadrant compared to the other three (first: $b = -0.06$, $SE = 0.01$, $p<0.001$,

193 second: $b = -0.08$, $SE = 0.01$, $p<0.001$, third: $b = -0.05$, $SE = 0.01$, $p<0.001$)

194 **Response Time and Accuracy**. Participants' response times and accuracies on

---

[1] `lme4` syntax: `lmer(proportion ~ quadrant + (1+quadrant|subject_id))`. Among other limitations, this approach may violate the independence assumptions of the linear model because looks to the four locations are not independent. This analysis was chosen because it is analogous to the ANOVA analysis conducted in the original paper.
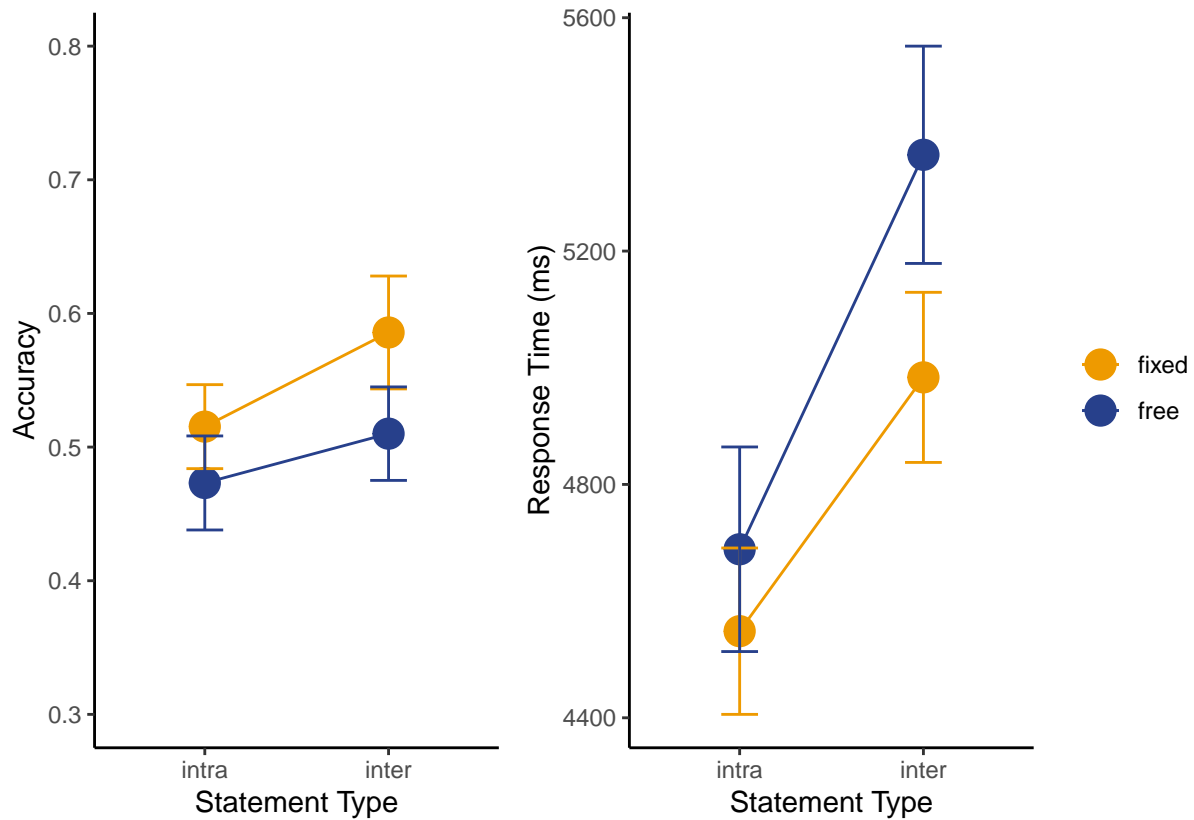
Figure 2. Accuracy and response times during memory retrieval.

memory questions are summarized in Figure 2. Both dependent variables were analyzed with linear mixed-effects model with relation type (interobject = -0.5, intraobject=0.5) and viewing_condition (fixed = -0.5, free=0.5) and their interaction as the predictors. The model included random intercepts for participants[2]. Response times were slower for interobject (e.g., "The train is to the right of the taxi.") than intraobject (e.g., "The train is facing right.") questions ($b = $ -555.60, $SE = $ 105.24, p<0.001). Response times were slower in the free viewing condition than the fixed condition ($b = $, $SE = $, p<0.001). Accuracy did not differ significantly between interobject and intraobject questions ($b = $ -0.05, $SE = $ 0.03, p=0.05). Participants were less accurate in the free viewing condition than the fixed condition ($b = $ -0.06, $SE = $ 0.03, p=0.03). The interaction ws not a

———

[2] lme4 syntax: `lmer(DV ~ relation_type*viewing_condition + (1|subject_id))`

signficant predictor for response times or accuracy.

Q: Can we recover item IDs to do crossed random effects

- here we will describe the analyses that are as close as possible to the original paper

- same but stricter validation cutoff

Calibration.

- here we will describe the analyses that correlate calibration quality with effect size at the individual level

Discussion

Experiment 3

The third study was a replication attempt of Manns, Stark, and Squire (2000) which aimed to show that the visual paired-comparison task, widely used in the patient literature, tapped into declarative memory. In the visual paired-comparison task, two identical pictures were presented side by side for a brief viewing period. After a delay, one of the previously viewed pictures was presented along with a new picture. Individuals looked more at the new picture than the old picture and the time spent looking was correlated with later recognition memory performance. On the other hand perceptual priming, thought to recruit non-declarative memory, was not linked to later recognition. (The perceptual priming arm of the design was not included in this replication.)

Methods

Participants. Our initial sample size was 51 participants for the first day of our experiment and 48 of them came back for the second day. Following Manns et al., we

excluded 3 participants due to perfect performance on the recognition memory test. Our final sample size was 45 participants.

Procedure. The task began with a 7-point eye-tracker calibration (each point was presented 3 times in a random order) and validation with 3 points (each presented once). The point locations were designed to focus calibration on the center of the screen and the middle of the left and right halves of the screen. The experiment was administered over the course of two consecutive days. It consisted of three sections: a presentation phase, a test phase, and a recognition test. The first two phases occurred on the first day, while the recognition test occurred on the second day.

During the presentation phase, participants viewed 24 pairs of identical color photographs depicting common objects. Each pair was presented for 5 seconds and an interval of 5 seconds elapsed before the next pair was shown. The order of the photographs was randomized and different for each participant. After completion of the presentation phase, participants were given a 5-minute break during which they could look away from the screen.

After the break, they were prompted to complete the eye-tracking calibration again before beginning the test phase. During this phase, participants again viewed 24 pairs of photographs with an interstimulus duration of 5 seconds. In each pair, one photograph was previously seen during the presentation phase, while the other was new. Which pictures were old or new was counterbalanced across participants. For half of the participants in each counterbalancing group, the new and old photographs were reversed.

Approximately 24 hours after completing the first session, with a leeway interval of 12 hours to accommodate busy schedules, participants were given the recognition test. It consisted of 48 photographs, presented one at a time. Each was shown on the screen for 1 second, followed by a 1 second interstimulus interval. Half of the pohotographs had been viewed twice on the previous day and were deemed the "targets." The other half depicted

252 an object with the same name as an object in one of the old photographs, but had not been

253 viewed before, deemed "foils." Each photograph remained on the screen until the

254 participants indicated whether or not they had seen it before by pressing 'y' for yes and 'n'

255 for no. After they pressed one of the two keys, a prompt on the screen asked them to rate

256 their confidence in their answer from 1 as a "pure guess" to 5 as "very sure." by clicking on

257 the corresponding number on the screen. No feedback on their responses was given during

258 the test.

259 The experimental design is visually depicted in Figure XX

260 **Materials.** Images were selected XXX...

261 There were two modifications we made to the methods of the original experiment. As

262 we are only replicating the declarative memory component of the original experiment, we

263 did not have a "priming group." Therefore, we followed only the procedure for the "looking

264 group." Additionally, for each section of the study, the stimuli was presented on a single

265 screen instead of two screens due to the constraints of the online experiment format.

266 **Data analysis.**

267 **Results**

268 **Discussion**

269 **Experiment 4**

270 The fourth study was a replication attempt of Experiment 1 in Ryskin, Qi, Duff, and

271 Brown-Schmidt (2017), which was closely modeled on Snedeker and Trueswell (2004).

272 These studies used the visual world paradigm to show that listeners use knowledge of the

273 co-occurrence statistics of verbs and syntactic structures to resolve ambiguity. For

274 example, in a sentence like "Feel the frog with the feather," the phrase "with the feather"

275 could be describing the frog, or it could be describing the instrument that should be used

276 to do the "feeling." When both options (a frog holding a feather and a feather by itself) are

277 available in the visual display, listeners rely on the verb's "bias" (statistical co-occurrence

278 either in norming or corpora) to rapidly choose an action while the sentence is unfolding. .

Methods

280   The stimuli, experimental code, and data and analysis scripts can be found on the

281 Open Science Framework at the following link, https://osf.io/x3c49/

282 (https://osf.io/x3c49/). The pre-registration for the study can be found at

283 https://osf.io/3v4pg (https://osf.io/3v4pg).

284   **Participants.**   58 (??) participants were paid $XX for their participation . A

285 sample size of 58 was chosen because we wanted to replicate the experiment with greater

286 statistical power. Note that the original study had a sample size of 24.

287   **Procedure.**

288   • *TO DO:* add details of calibration point locations

289   After the eye-tracking calibration, participants went through an audio test so they

290 could adjust the audio on their computer to a comfortable level. Before beginning the

291 experiment, they were given instructions that four objects would appear, an audio prompt

292 would play, and they should do their best to use their mouse to act out the instructions.

293 They then went through three practice trials which were followed by 54 critical trials and

294 24 filler trials presented in a random order.

295   During a trial, four pictures were displayed (target animal, target instrument,

296 distractor animal, distractor instrument), one in each corner of the screen, and participants

297 heard an audio prompt that contained instructions about the action they needed to act out

298 (e.g., "Rub the butterfly with the crayon"; see Figure XX)[3]. Using their cursor,

———

[3] In the original study, the pictures appeared one by one on the screen and their names were played as they

299 participants could act out the instructions by clicking on objects and moving them or

300 motioning over the objects[4]. After the action was completed, the participants were

301 instructed to press the space bar which led to a screen that said "Click Here" in the middle

302 in order to remove bias in the eye and mouse movements from the previous trial. The

303 experiment only allowed the participants to move on to the next trial once the audio was

304 completely done playing and the mouse had been moved over at least one object.

305   *TO DO*: ADD FIGURES Figure 1: An example of a critical trial for the sentence

306 "Rub the butterfly with the crayon." The butterfly is the target animal, the panda is the

307 distractor animal, the crayon is the target instrument, and the violin is the distractor

308 instrument.

309   **Materials.**   The images and audios presented to the participants were the same

310 stimuli used in the original study (available here). The critical trials were divided into

311 modifier-biased, instrument-biased, and equibiased conditions, and the filler trials did not

312 contain ambiguous instructions. Two lists of critical trials were made with different verb

313 and instrument combinations (e.g., "rub" could be paired with "panda" and "crayon" in

314 one list and "panda" and "violin" in the second list). Within each list, the same verb was

315 presented twice but each time with a different target instrument and animal. The lists were

316 randomly assigned to the participants to make sure the effects were not caused by the

317 properties of the animal or instrument images used. The list of verbs used can be found in

318 Appendix A of the original study.

---

appeared. We removed this introductory portion of the trial to save time

[4] As opposed to the original study we recorded mouse movement instead of clicking behavior since not all
of the audio prompts required clicking. For example, the sentence "locate the camel with the straw" may
not involve any clicking but rather only mousing over the camel.

Results

320   **Replication.** The location of initial mouse movements was used to assess whether
321 the final interpretation of ambiguous sentences was biased by the verb. Figure 3 suggests
322 that listeners were more likely to move their mouse first over the target instrument when
323 the verb was equi-biased than when the verb was modifier-biased and even more so when
324 the verb was instrument-biased. The opposite graded pattern can be observed for mouse
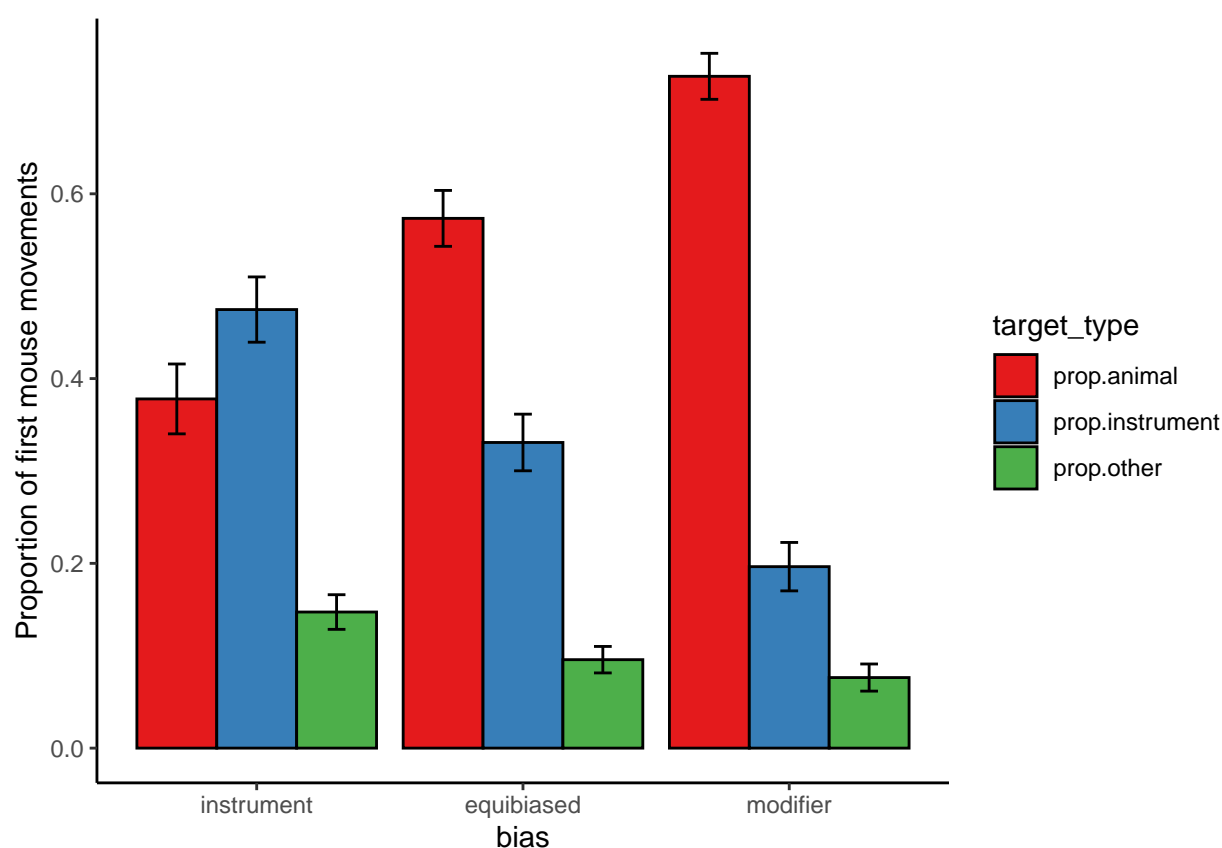325 movements over the target animal.

Figure 3. Proportion of first mouse movements by location and verb bias.

326   A mixed-effects logistic regression model was used to predict whether the first
327 movement was on the target instrument with the verb bias condition as an orthogonally
328 contrast-coded (instrument vs. equi & modifier: inst = -2/3, equi = 1/3, mod = 1/3; equi
329 vs. modifier: inst = 0, equi = -1/2, mod = 1/2) fixed effect. Participants and items were

entered as varying intercepts with by-participant varying slopes for verb bias condition[5].

Participants were more likely to first move their mouse over target instruments in the

instrument-biased condition relative to the equi-biased and modifier-biased condition ($b =$

-1.50, $SE = 0.25$, $p < 0.01$). Further, participants were more likely to first move their

mouse over target instruments in the equi-biased condition relative to the modifier-biased

condition ($b =$ -1.10, $SE = 0.29$, $p < 0.01$)

Gaze fixations were time-locked to the auditory stimulus on a trial by trial basis and

categorized as being directed towards one of the four items in the display if the x, y

coordinates fell within a rectangle containing the image. Figure 4 suggests that the

participants made more fixations to the target animal when the verb was modifier-biased...

In order to assess how verb bias impacted sentence disambiguation as the sentence

unfolded, the proportion of fixations was computed in three time windows: the

verb-to-animal window (from verb onset + 200 ms to animal onset + 200 ms), the

animal-to-instrument window (from animal onset + 200 ms to instrument onset + 200 ms),

and the post-instrument window (from instrument onset + 200 ms to instrument onset +

1500ms + 200 ms). Mixed-effects linear regression models were used to predict the

proportions of fixations to the target animal within each time window with the verb bias

condition as an orthogonally contrast-coded (instrument vs. equi & modifier: inst = -2/3,

equi = 1/3, mod = 1/3; equi vs. modifier: inst = 0, equi = -1/2, mod = 1/2) fixed effect.

Participants and items were entered as varying intercepts[6]. In the *verb-to-noun* window,

participants did not look more at the target animal in any of the verb bias conditions

(Instrument vs. Equi and Modifier: $b =$ -0.01, $SE = 0.02$, $p = 0.59$; Equi vs. Modifier: $b =$

---

[5] lme4 syntax: `glmer(is.mouse.over.instrument ~ verb_bias + (1 + verb_bias | participant) + (1 | item), family="binomial", data=d)`

[6] lme4 syntax: `lmer(prop.fix.target.animal ~ verb_bias + (1 + verb_bias | participant) + (1 | item), data=d)`. A model with by-participant varying slopes for verb bias condition was first attempted but did not converge.
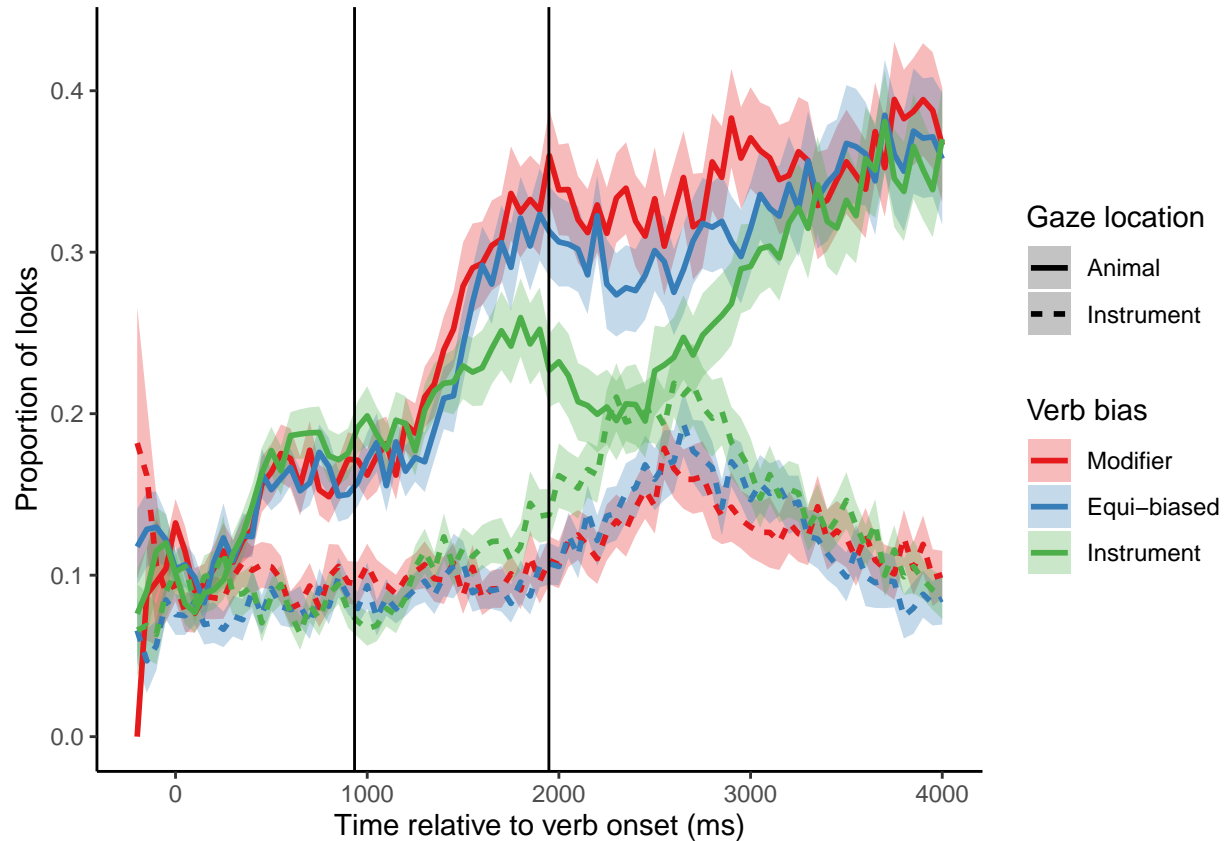
Figure 4. Timecourse of eye-gaze to target animal and target instrument by verb bias condition. Vertical lines indicate average onsets of animal and instrument offset by 200ms.

352   0, $SE = 0.02$, $p = 1$). In the *noun-to-instrument* window, participants looked more at the

353   target animal in the modifier-biased condition relative to the equi-biased and

354   instrument-biased condition ($b = 0.03$, $SE = 0.01$, $p < 0.01$) and in the equi-biased

355   relative to the instrument-biased condition ($b = 0.02$, $SE = 0.01$, $p < 0.05$). In the

356   *post-instrument* window, participants looked more at the target animal in the

357   modifier-biased condition relative to the equi-biased and instrument-biased condition ($b =$

358   0.08, $SE = 0.02$, $p < 0.01$) but not significantly so in the equi-biased relative to the

359   instrument-biased condition ($b = 0.03$, $SE = 0.02$, $p = 0.59$).

360   **Comparison to in-lab data.**

361   • here we will describe a direct comparison to data collected in the lab

Calibration.

- here we will describe the analyses that correlate calibration quality with effect size at the level of the individual

Effects of ROIs.

- here we will describe how results change depending on the size of the ROIs (using the image vs the screen quadrant)

Discussion

<div align="center">

Experiment 5

</div>

The fifth study was a replication attempt of Shimojo, Simion, Shimojo, and Scheier (2003), which found that human gaze is actively involved in preference formation. Participants were shown pairs of human faces and were asked to choose the face which they found more attractive. Prior to making their explicit selection, participants spent more time looking at the face that was ultimately judged more attractive. This bias was not as large when the task was face shape discrimination rather than a preference-based choice.

Methods

All stimuli, experiment scripts, data, and analysis scripts are available on the Open Science Framework at https://osf.io/eubsc/ (https://osf.io/eubsc/). The study pre-registration is available at https://osf.io/tv57s (https://osf.io/tv57s).

Participants.   27 participants were recruited on Prolific to participate in stimulus norming (for attractiveness). They were paid $XX for completing the experiment. Data from 3 participants was excluded because their mode response made up more than 50% of their total responses , for a total of 24 participants in the norming.

384   50 participants for the main task were recruited on Prolific and were paid $XX. 8

385 subjects, 4 from the attractiveness task group and 4 from the roundness task group, were

386 excluded for incorrect validations. After this data exclusion, we ended up with 21

387 participants each for the attractiveness task and the roundness task.The original sample

388 size in Shimojo et al. (2003) was 10 participants total.

389   **Procedure and Design.**   At the beginning of the experimental task, participants

390 completed a 9-point eye-tracker calibration (each point appeared 3 times in random order)

391 and 3-point validation. The validation point appeared once at center, middle left, and

392 middle right locations in random order.

393   During each trial of the main task, two faces were displayed on the two halves of the

394 screen, one on the left and one on the right (as in Figure XX). Participants were randomly

395 assigned to one of two tasks: attractiveness or shape judgment. In the attractiveness task,

396 participants were asked to chose the more attractice face in the pair and in the shape

397 judgment task participants were asked to pick the face that appeared rounder. They

398 pressed the "a" key on their keyboard to select the face on the left and the "d" key to select

399 the face on the right. A fixation cross appeared in the center of the screen between each set

400 of faces. Participants were asked to look at this fixation cross in order to reset their gaze in

401 between trials (???). The order of the 19 face pairs was random for each participant.

402   **Materials and Norming.**   The faces in our replication were selected from a set of

403 1,000 faces within the Flickr-Faces-HQ Dataset. (The face images used in Shimojo et

404 al. were from the Ekman face database and the AR face database.) These images were

405 chosen because the person in each image was looking at the camera with a fairly neutral

406 facial expression and appeared to be over the age of 18. Participants in the norming study

407 viewed 172 faces and were asked to rate them on a scale from 1 (less attractive) to 7 (more

408 attractive) using a slider. Faces were presented one at a time and in a random order for

409 each participant. Following Shimojo et al., 19 face pairs were made by matching two faces

410 that had a difference in mean attractiveness ratings that was 0.25 points or lower and that

matched in gender, race, and age group (young adult, adult, or older adult).

Data analysis. In the original study, a video-based eye tracker was used. The eye movements of participants were recorded with a digital camera, and eye position was then tracked using MediaAnalyzer software. This eye tracker was able to collect 30 samples per second. In our study, the camera on the device used by the participant and the jsPsych WebGazer package were used to determine where a participant's gaze was directed. Since our eye tracking method did not have the same sample collection rate as the eye tracker in Shimojo et al., time windows of 50 milliseconds were used when looking at the likelihood of gazing at the chosen face at a given point between a cutoff time and when the decision was made.

Results

Discussion

Experiment 6

The sixth study was a replication attempt of Posner et al. ??.

Methods

Participants.

Procedure and Design.

Data analysis.

Results

Replication.

Calibration.

ROIs.

Item Numbers.

Discussion

Combined Analyses

- Pooling data from all experiments we can look at patterns in the calibration and validation data

General Discussion

References

Altmann, G. T. M., & Kamide, Y. (1999). Incremental interpretation at verbs: Restricting the domain of subsequent reference. *Cognition*, *73*(3), 247–264. https://doi.org/10.1016/S0010-0277(99)00059-1

Aust, F., & Barth, M. (2020). *papaja: Create APA manuscripts with R Markdown.* Retrieved from https://github.com/crsh/papaja

Bates, D., Mächler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, *67*(1), 1–48. https://doi.org/10.18637/jss.v067.i01

Bates, D., & Maechler, M. (2021). *Matrix: Sparse and dense matrix classes and methods.* Retrieved from https://CRAN.R-project.org/package=Matrix

Bolker, B., & Robinson, D. (2020). *Broom.mixed: Tidying methods for mixed models.* Retrieved from https://CRAN.R-project.org/package=broom.mixed

Chang, W., Cheng, J., Allaire, J., Sievert, C., Schloerke, B., Xie, Y., . . . Borges, B. (2021). *Shiny: Web application framework for r.* Retrieved from https://CRAN.R-project.org/package=shiny

de Leeuw, J. R. (2015). jsPsych: A JavaScript library for creating behavioral experiments in a Web browser. *Behavior Research Methods*, *47*(1), 1–12. https://doi.org/10.3758/s13428-014-0458-y

Johansson, R., & Johansson, M. (2014). Look Here, Eye Movements Play a Functional Role in Memory Retrieval. *Psychological Science*, *25*(1), 236–242. https://doi.org/10.1177/0956797613498260

Kuznetsova, A., Brockhoff, P. B., & Christensen, R. H. B. (2017). lmerTest package: Tests in linear mixed effects models. *Journal of Statistical Software*, *82*(13), 1–26. https://doi.org/10.18637/jss.v082.i13

Manns, J. R., Stark, C. E. L., & Squire, L. R. (2000). The visual paired-comparison task as a measure of declarative memory. *Proceedings of the National Academy*

of Sciences, *97*(22), 12375–12379. https://doi.org/10.1073/pnas.220398097

Ooms, J. (2014). The jsonlite package: A practical and consistent mapping between JSON data and r objects. *arXiv:1403.2805 [Stat.CO]*. Retrieved from https://arxiv.org/abs/1403.2805

Papoutsaki, A., Sangkloy, P., Laskey, J., Daskalova, N., Huang, J., & Hays, J. (2016). WebGazer: Scalable webcam eye tracking using user interactions. In *Proceedings of the 25th international joint conference on artificial intelligence (IJCAI)* (pp. 3839–3845). AAAI.

R Core Team. (2021). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from https://www.R-project.org/

Ryskin, R., Qi, Z., Duff, M. C., & Brown-Schmidt, S. (2017). Verb biases are shaped through lifelong learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, *43*(5), 781–794. https://doi.org/10.1037/xlm0000341

Shimojo, S., Simion, C., Shimojo, E., & Scheier, C. (2003). Gaze bias both reflects and influences preference. *Nature Neuroscience*, *6*(12), 1317–1322. https://doi.org/10.1038/nn1150

Singmann, H., Bolker, B., Westfall, J., Aust, F., & Ben-Shachar, M. S. (2021). *Afex: Analysis of factorial experiments*. Retrieved from https://CRAN.R-project.org/package=afex

Snedeker, J., & Trueswell, J. C. (2004). The developing constraints on parsing decisions: The role of lexical-biases and referential scenes in child and adult sentence processing. *Cognitive Psychology*, *49*(3), 238–299. https://doi.org/10.1016/j.cogpsych.2004.03.001

Spivey, M. J., & Geng, J. J. (2001). Oculomotor mechanisms activated by imagery and memory: Eye movements to absent objects. *Psychological Research*, *65*(4), 235–241. https://doi.org/10.1007/s004260100059

Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science*, *268*(5217), 1632–1634.

Wickham, H. (2016). *ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York. Retrieved from https://ggplot2.tidyverse.org

Wickham, H. (2019). *Stringr: Simple, consistent wrappers for common string operations*. Retrieved from https://CRAN.R-project.org/package=stringr

Wickham, H. (2021a). *Forcats: Tools for working with categorical variables (factors)*. Retrieved from https://CRAN.R-project.org/package=forcats

Wickham, H. (2021b). *Tidyr: Tidy messy data*. Retrieved from https://CRAN.R-project.org/package=tidyr

Wickham, H., François, R., Henry, L., & Müller, K. (2021). *Dplyr: A grammar of data manipulation*. Retrieved from https://CRAN.R-project.org/package=dplyr

Wickham, H., & Hester, J. (2020). *Readr: Read rectangular text data*. Retrieved from https://CRAN.R-project.org/package=readr